

Generative AI and RAG – Demo Document

1. What is Generative AI?

Generative AI is a type of artificial intelligence that can create new content such as text, images, audio, and code. It learns patterns from large amounts of data and generates content based on a user's prompt.

2. What is an LLM?

LLM stands for Large Language Model. An LLM processes text and generates language-based responses. Examples of tasks include answering questions, summarizing text, translating languages, and generating code.

3. What are Tokens?

A token is a small unit of text processed by a language model. A token can be a whole word, part of a word, or sometimes punctuation. LLM context windows are measured in tokens.

4. What are Embeddings?

Embeddings are numerical vector representations of text. Texts with similar meanings tend to have similar vector representations. Embeddings are commonly used for semantic search and retrieval.

5. What is a Vector Database?

A vector database stores embeddings and supports similarity search. In a RAG application, it can be used to find document chunks that are semantically relevant to a user's question.

6. What is RAG?

RAG stands for Retrieval-Augmented Generation. A RAG system retrieves relevant information from a knowledge source and provides that information as context to an LLM before generating an answer.

7. Basic RAG Flow

Document → Chunking → Embeddings → Vector Database → User Question → Retrieval → Relevant Context → LLM → Answer

8. Why use RAG?

RAG can help an application answer questions using information from a specific document or knowledge base. It can also reduce unsupported answers by instructing the LLM to use retrieved context.